

Introduction to R

Matthew Suderman

**Lecturer in Epigenetic
Epidemiology**



Things you can do with R

| Description | R tool |
|--|-----------------------------|
| Remember data | Variables |
| Calculate | Maths functions |
| Store and manipulate sequences of values | Vectors and lists |
| Calculate with vectors | Vector operations |
| Generate data | Vector-generating functions |
| Store and manipulate matrices | Matrices |
| Calculate with matrices | Matrix functions |
| Store and manipulate datasets | Data frames |
| Handle missing data | NA |
| Save and load data | Files |
| Statistical analyses | Statistics functions |
| Visualise data | Plots |
| Create and apply recipes | Functions |
| Repeatedly apply recipes | Loops and apply functions |
| Make decisions | If/else statements |
| Share | Packages |

Getting started

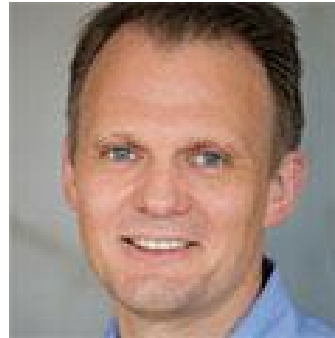
what, why, downloading, opening, using, rstudio, help, errors

Getting started: what

- R is a statistical programming language (based on S)
- R is open source - researchers develop packages to implement new statistical methods, plots or applications
- R runs on Windows, MacOS and UNIX



John Chambers



Dirk Eddelbuettel

Getting started: why

... rather than S, STATA, SPSS, etc

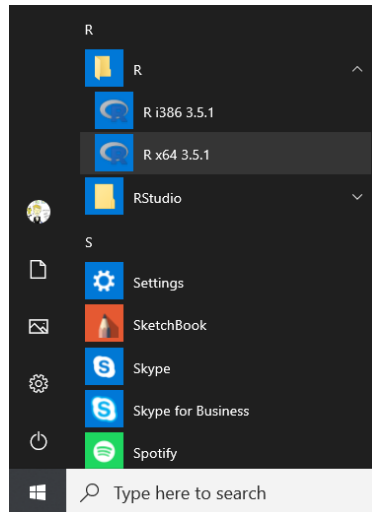
- R is free!
- R is flexible
- R is good at handling large datasets and multiple objects
- R has good plotting tools and packages for statistical analysis

Getting started: installing

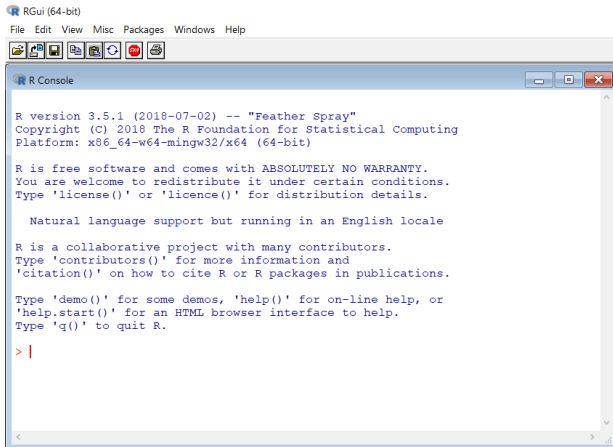
- Visit <https://www.r-project.org> and click "Download R"
- Choose your nearest CRAN mirror (such as <http://www.stats.bris.ac.uk/R>)
- Choose "Download R for [Windows/Mac/Linux]"
- Choose "base" for Windows and click on "Download R 3.5.2 for Windows"
- Choose "R-3.5.2.pkg" for Mac
- Once the .exe (Windows) or .pkg (Mac) have downloaded, run and install

Getting started: opening

- Click "Start" | "All Programs" | "R" | "R x64 3.6.0"



Getting started: navigating



```
RGui (64-bit)
File Edit View Misc Packages Windows Help

R Console

R version 3.5.1 (2018-07-02) -- "Feather Spray"
Copyright (C) 2018 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

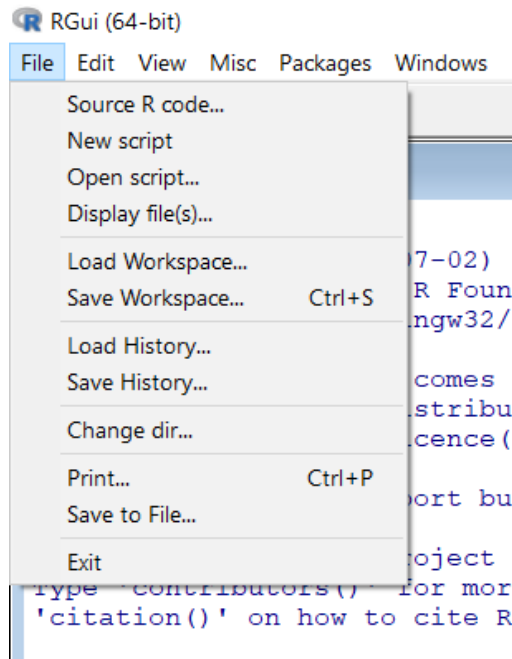
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> |
```

- **Drop-down menus** are available but R cannot be used in a point-and-click fashion
- The **R console** initially just shows information on the version of R running, how to get help and how to quit
- At the bottom of the console is the **R command line** where you can tell R what to do

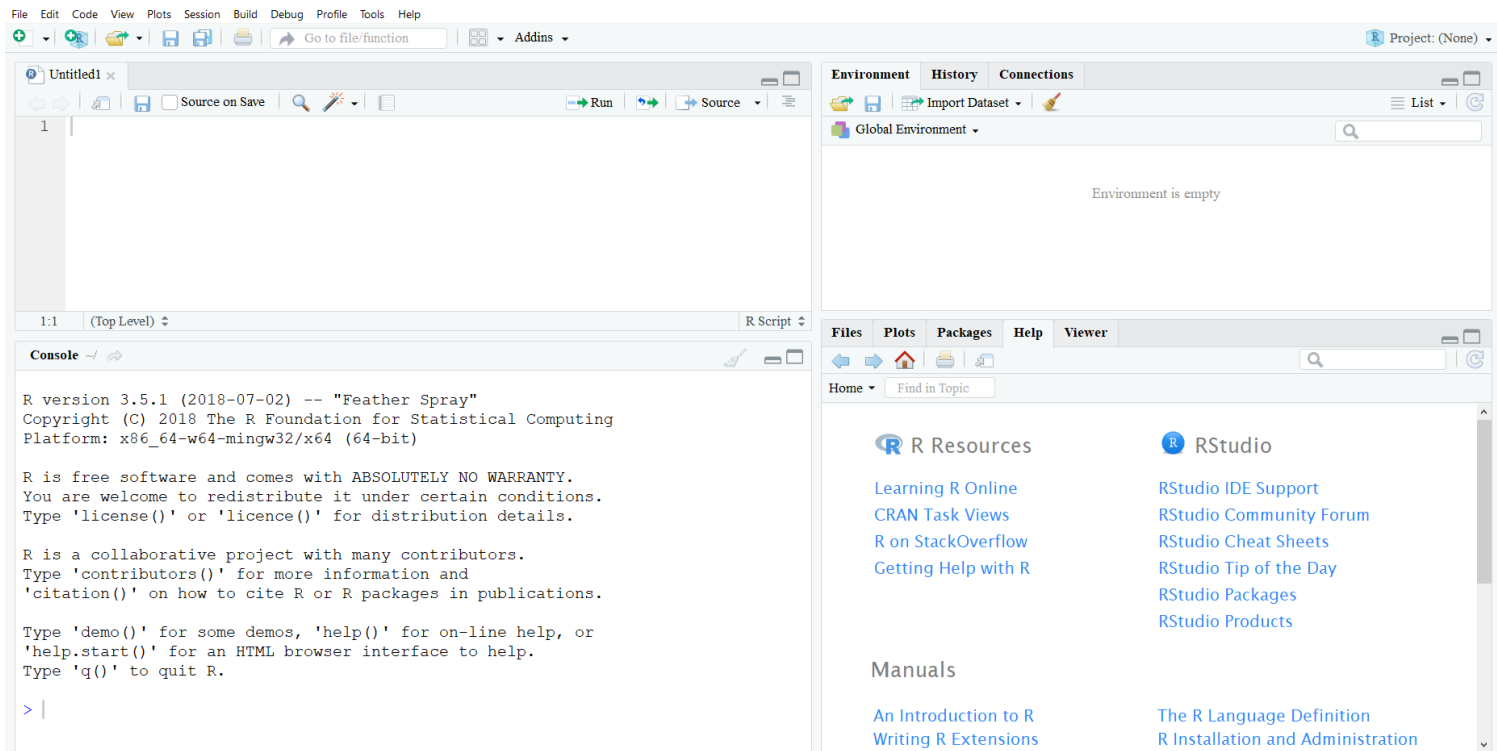
Getting started: menus



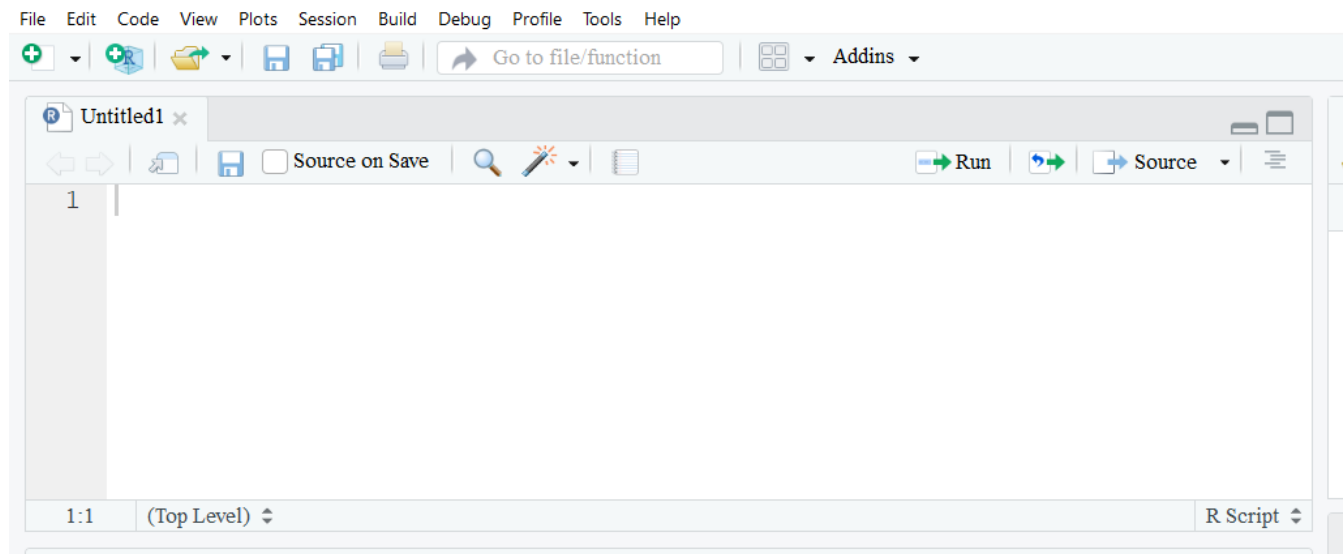
- A **script** is sequence of R commands. Creating or opening a script will open the script in a text editor.
- A **workspace** includes all the objects which are in R's memory at a given time. These can include data or results. To see what is in the current workspace type `ls()`.
- The **history** includes all the commands which are displayed during an R session

Getting started: RStudio

- RStudio makes R easier to use
- It includes a code editor, debugging & visualization tools
- <https://www.rstudio.com>



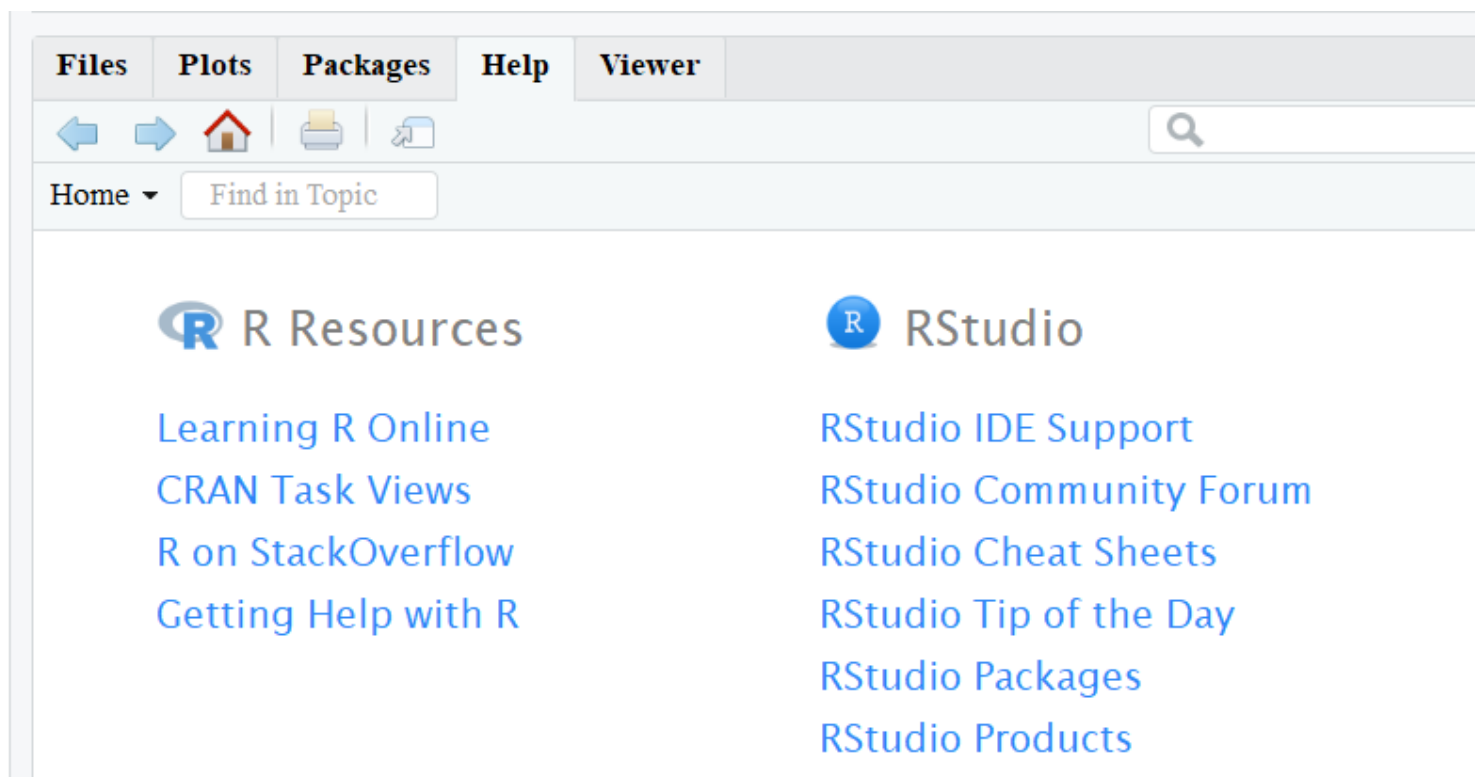
Getting started: RStudio script window



In the script window, you can **save a script** (disk image), **run** selected commands ("Run") and **run** the entire script ("Script").

Getting started: help!

- Use the command `help(topic)` or `?topic`
- In R this will open a web browser
- In RStudio this will open the Help tab



Getting started: `help(mean)`

Getting started: `help(mean)`

`mean {base}`

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

- `x` An R object. Currently there are methods for numeric/logical vectors and date, date-time and time interval objects. Complex vectors are allowed for `trim = 0`, only.
- `trim` the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- `na.rm` a logical value indicating whether NA values should be stripped before the computation proceeds.
- `...` further arguments passed to or from other methods.

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

- `x` An R object. Currently there are methods for numeric/logical vectors and date, date-time and time interval objects. Complex vectors are allowed for `trim = 0`, only.
- `trim` the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- `na.rm` a logical value indicating whether NA values should be stripped before the computation proceeds.
- `...` further arguments passed to or from other methods.

Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

- `x` An R object. Currently there are methods for numeric/logical vectors and date, date-time and time interval objects. Complex vectors are allowed for `trim = 0`, only.
- `trim` the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- `na.rm` a logical value indicating whether NA values should be stripped before the computation proceeds.
- `...` further arguments passed to or from other methods.

Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) The New S Language. Wadsworth & Brooks/Cole.

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

- `x` An R object. Currently there are methods for numeric/logical vectors and date, date-time and time interval objects. Complex vectors are allowed for `trim = 0`, only.
- `trim` the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- `na.rm` a logical value indicating whether NA values should be stripped before the computation proceeds.
- `...` further arguments passed to or from other methods.

Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) The New S Language. Wadsworth & Brooks/Cole.

See Also

`weighted.mean`, `mean.POSIXct`, `colMeans` for row and column means.

Getting started: `help(mean)`

```
mean {base}
```

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

- `x` An R object. Currently there are methods for numeric/logical vectors and date, date-time and time interval objects. Complex vectors are allowed for `trim = 0`, only.
- `trim` the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- `na.rm` a logical value indicating whether NA values should be stripped before the computation proceeds.
- `...` further arguments passed to or from other methods.

Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) The New S Language. Wadsworth & Brooks/Cole.

See Also

`weighted.mean`, `mean.POSIXct`, `colMeans` for row and column means.

Examples

```
x <- c(0:10, 50)
xm <- mean(x)
c(xm, mean(x, trim = 0.10))
```

Getting started: error messages

Suppose you type the following but variable 'a' has not been created.

```
> 10*a
```

Getting started: error messages

Suppose you type the following but variable 'a' has not been created.

```
> 10*a
```

```
Error: object 'a' not found
```

Getting started: error messages

Suppose you type the following but variable 'a' has not been created.

```
> 10*a
```

```
Error: object 'a' not found
```

R provides informative error messages

Getting started: error messages

Suppose you type the following but variable 'a' has not been created.

```
> 10*a
```

```
Error: object 'a' not found
```

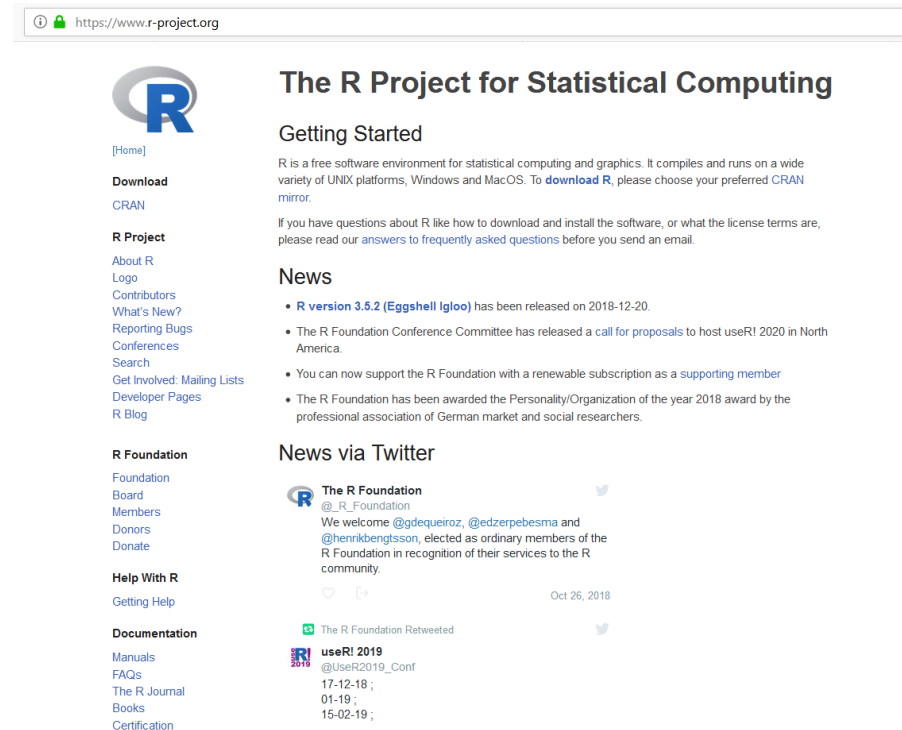
R provides informative error messages



Getting started: www.r-project.org

Note:

- R Project | Search
- Help with R
- Manuals
- FAQs



The screenshot shows the homepage of the R Project for Statistical Computing. The browser address bar displays <https://www.r-project.org>. The page features the R logo, a navigation menu on the left, and a main content area with sections for 'Getting Started', 'News', and 'News via Twitter'.

The R Project for Statistical Computing

Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To [download R](#), please choose your preferred [CRAN mirror](#).

If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

News

- **R version 3.6.2 (Eggshell Igloo)** has been released on 2018-12-20.
- The R Foundation Conference Committee has released a [call for proposals](#) to host useR! 2020 in North America.
- You can now support the R Foundation with a renewable subscription as a [supporting member](#)
- The R Foundation has been awarded the Personality/Organization of the year 2018 award by the professional association of German market and social researchers.

News via Twitter

The R Foundation @_R_Foundation

We welcome [@gdequeiroz](#), [@edzerpebesma](#) and [@henrikbengtsson](#), elected as ordinary members of the R Foundation in recognition of their services to the R community.

Oct 26, 2018

The R Foundation Retweeted

useR! 2019 @UseR2019_Conf

17-12-18 ;
01-19 ;
15-02-19 ;

Getting started: forums

- Mailing list archive and forum <http://r.789695.n4.nabble.com>
- Stack Overflow <https://stackoverflow.com>
- Google!

Things you can do with R

*remember, calculate, store, generate, analyse, visualise, save, repeat, apply,
decide, share*

Remember data

A **variable** is a storage location for data with a name.

Store data ...

```
> x <- 3
> y <- 4.2
> name <- "abcdef"
> skip <- TRUE
```

Recall data ...

```
> x
[1] 3
> name
[1] "abcdef"
> nchar(name)
[1] 6
> z
Error: object 'z' not found
```

Query data type ...

```
> is.numeric(x)
[1] TRUE
> is.logical(name)
[1] FALSE
```

Change data type ...

```
> as.character(x)
[1] "3"
> as.numeric("3.14")
[1] 3.14
> as.logical("abc")
[1] NA
> as.logical(x)
[1] TRUE
```

Calculate

```
> x + y + 0.8  
[1] 8
```

Calculate

```
> x + y + 0.8  
[1] 8
```

```
> 4*x/2  
[1] 6
```

Calculate

```
> x + y + 0.8  
[1] 8
```

```
> 4*x/2  
[1] 6
```

```
> log10(100)  
[1] 100  
> 2^x  
[1] 8  
> sin(pi/2)  
[1] 1  
> ln(exp(1))  
[1] 1  
> sqrt(x^2)  
[1] 3
```

Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```


Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```

```
[1] 11 12 13 14 15 16 17 18 19 20
```

Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```

```
[1] 11 12 13 14 15 16 17 18 19 20
```

```
> x[3] ## show the third
```

Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```

```
[1] 11 12 13 14 15 16 17 18 19 20
```

```
> x[3] ## show the third
```

```
[1] 13
```

Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```

```
[1] 11 12 13 14 15 16 17 18 19 20
```

```
> x[3] ## show the third
```

```
[1] 13
```

```
> length(x)
```

Store and manipulate sequences of values

A **vector** is a sequence of variables of the same type.

```
> x <- c(11,12,13,14,15,16,17,18,19,20)
> x
```

```
[1] 11 12 13 14 15 16 17 18 19 20
```

```
> x[3] ## show the third
```

```
[1] 13
```

```
> length(x)
```

```
[1] 10
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```


Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```

```
[1] 12 13 14 16 18 19 20
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```

```
[1] 12 13 14 16 18 19 20
```

```
> x[x>13]
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```

```
[1] 12 13 14 16 18 19 20
```

```
> x[x>13]
```

```
[1] 14 15 16 17 18 19 20
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```

```
[1] 12 13 14 16 18 19 20
```

```
> x[x>13]
```

```
[1] 14 15 16 17 18 19 20
```

```
y <- c("Dog", "Cat")  
z <- c(x[1:3], y)  
z
```

Store and manipulate sequences of values (continued)

Vectors can be subset and combined.

```
> x[c(1, 5, 7)]
```

```
[1] 11 15 17
```

```
> x[-c(1, 5, 7)]
```

```
[1] 12 13 14 16 18 19 20
```

```
> x[x>13]
```

```
[1] 14 15 16 17 18 19 20
```

```
y <- c("Dog", "Cat")  
z <- c(x[1:3], y)  
z
```

```
[1] "11" "12" "13" "Dog" "Cat"
```

Calculate with vectors

```
> x+1
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```


Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```

```
[1] 33 36 39 42 45 48 51 54 57 60
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```

```
[1] 33 36 39 42 45 48 51 54 57 60
```

```
> x/2
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```

```
[1] 33 36 39 42 45 48 51 54 57 60
```

```
> x/2
```

```
[1] 5.5 6 6.5 7 7.5 8 8.5 9 9.5 10
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```

```
[1] 33 36 39 42 45 48 51 54 57 60
```

```
> x/2
```

```
[1] 5.5 6 6.5 7 7.5 8 8.5 9 9.5 10
```

```
> log10(x)
```

Calculate with vectors

```
> x+1
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
> 3*x
```

```
[1] 33 36 39 42 45 48 51 54 57 60
```

```
> x/2
```

```
[1] 5.5 6 6.5 7 7.5 8 8.5 9 9.5 10
```

```
> log10(x)
```

```
[1] 1.041393 1.079181 1.113943 1.146128 1.176091 1.204120 1.230449 1.255273  
[9] 1.278754 1.301030
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```


Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```

```
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```

```
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
```

```
> x*y
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```

```
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
```

```
> x*y
```

```
[1] 132 156 182 210 240 272 306 342 380 420
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```

```
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
```

```
> x*y
```

```
[1] 132 156 182 210 240 272 306 342 380 420
```

```
> x/(y-1)
```

Calculate with vectors (continued)

```
> y <- x + 1  
> x + y
```

```
[1] 23 25 27 29 31 33 35 37 39 41
```

```
> x-y
```

```
[1] -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
```

```
> x*y
```

```
[1] 132 156 182 210 240 272 306 342 380 420
```

```
> x/(y-1)
```

```
[1] 1 1 1 1 1 1 1 1 1 1
```

Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

```
[1] 0 5 10 15 20 25 30 35 40 45 50  
[12] 55 60 65 70 75 80 85 90 95 100
```

Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

```
[1] 0 5 10 15 20 25 30 35 40 45 50  
[12] 55 60 65 70 75 80 85 90 95 100
```

```
> length(a)
```


Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

```
[1] 0 5 10 15 20 25 30 35 40 45 50  
[12] 55 60 65 70 75 80 85 90 95 100
```

```
> length(a)
```

```
[1] 21
```

Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

```
[1] 0 5 10 15 20 25 30 35 40 45 50  
[12] 55 60 65 70 75 80 85 90 95 100
```

```
> length(a)
```

```
[1] 21
```

```
> a <- seq(0, 100, length=5)  
> a
```

Generate data: seq()

```
> a <- seq(0, 100, by=5)  
> a
```

```
[1] 0 5 10 15 20 25 30 35 40 45 50  
[12] 55 60 65 70 75 80 85 90 95 100
```

```
> length(a)
```

```
[1] 21
```

```
> a <- seq(0, 100, length=5)  
> a
```

```
[1] 0 25 50 75 100
```

Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)
> b
```

Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)  
> b
```

```
[1] 1 1 1 1 1
```

Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)  
> b
```

```
[1] 1 1 1 1 1
```

```
> b <- rep(1, 5)  
> b
```

Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)  
> b
```

```
[1] 1 1 1 1 1
```

```
> b <- rep(1, 5)  
> b
```

```
[1] 1 1 1 1 1
```

Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)
> b
```

```
[1] 1 1 1 1 1
```

```
> b <- rep(1, 5)
> b
```

```
[1] 1 1 1 1 1
```

```
> rep(c(1,2,3),4)
```


Generate data: `rep()`

```
> b <- c(1, 1, 1, 1, 1)  
> b
```

```
[1] 1 1 1 1 1
```

```
> b <- rep(1, 5)  
> b
```

```
[1] 1 1 1 1 1
```

```
> rep(c(1,2,3),4)
```

```
[1] 1 2 3 1 2 3 1 2 3 1 2 3
```

Store and manipulate matrices

A **matrix** is a set of elements laid out in rows and columns.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> z <- matrix(c(x, y), nrow=5)
> z
```

Store and manipulate matrices

A **matrix** is a set of elements laid out in rows and columns.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> z <- matrix(c(x, y), nrow=5)
> z
```

```
[,1] [,2]
[1,] "1" "A"
[2,] "1" "B"
[3,] "1" "A"
[4,] "1" "B"
[5,] "1" "C"
```

Store and manipulate matrices

A **matrix** is a set of elements laid out in rows and columns.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> z <- matrix(c(x, y), nrow=5)
> z
```

```
[,1] [,2]
[1,] "1" "A"
[2,] "1" "B"
[3,] "1" "A"
[4,] "1" "B"
[5,] "1" "C"
```

```
> dim(z) ## dimensions
```

Store and manipulate matrices

A **matrix** is a set of elements laid out in rows and columns.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> z <- matrix(c(x, y), nrow=5)
> z
```

```
[,1] [,2]
[1,] "1" "A"
[2,] "1" "B"
[3,] "1" "A"
[4,] "1" "B"
[5,] "1" "C"
```

```
> dim(z) ## dimensions
```

```
[1] 5 2
```

Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

```
[,1] [,2] [,3] [,4] [,5]  
[1,] "1" "1" "1" "1" "1"  
[2,] "A" "B" "A" "B" "C"
```

Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

```
[,1] [,2] [,3] [,4] [,5]  
[1,] "1" "1" "1" "1" "1"  
[2,] "A" "B" "A" "B" "C"
```

```
> cbind(0:4, z)
```


Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

```
[,1] [,2] [,3] [,4] [,5]  
[1,] "1" "1" "1" "1" "1"  
[2,] "A" "B" "A" "B" "C"
```

```
> cbind(0:4, z)
```

```
[1,] "0" "1" "A"  
[2,] "1" "1" "B"  
[3,] "2" "1" "A"  
[4,] "3" "1" "B"  
[5,] "4" "1" "C"
```

Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

```
[,1] [,2] [,3] [,4] [,5]  
[1,] "1" "1" "1" "1" "1"  
[2,] "A" "B" "A" "B" "C"
```

```
> cbind(0:4, z)
```

```
[1,] "0" "1" "A"  
[2,] "1" "1" "B"  
[3,] "2" "1" "A"  
[4,] "3" "1" "B"  
[5,] "4" "1" "C"
```

```
> rbind(c("0", "X"), z)
```

Store and manipulate matrices (continued)

```
> t(z) ## transpose
```

```
[,1] [,2] [,3] [,4] [,5]  
[1,] "1" "1" "1" "1" "1"  
[2,] "A" "B" "A" "B" "C"
```

```
> cbind(0:4, z)
```

```
[1,] "0" "1" "A"  
[2,] "1" "1" "B"  
[3,] "2" "1" "A"  
[4,] "3" "1" "B"  
[5,] "4" "1" "C"
```

```
> rbind(c("0", "X"), z)
```

```
[1,] "0" "X"  
[2,] "1" "A"  
[3,] "1" "B"  
[4,] "1" "A"  
[5,] "1" "B"  
[6,] "1" "C"
```

Store and manipulate matrices (continued)

```
> z[,1]
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

```
[1] "A"
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

```
[1] "A"
```

```
> z[2,]
```


Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

```
[1] "A"
```

```
> z[2,]
```

```
[1] "1" "B"
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

```
[1] "A"
```

```
> z[2,]
```

```
[1] "1" "B"
```

```
> z[c(1,3),2]
```

Store and manipulate matrices (continued)

```
> z[,1]
```

```
[1] "1" "1" "1" "1" "1"
```

```
> z[1,2]
```

```
[1] "A"
```

```
> z[2,]
```

```
[1] "1" "B"
```

```
> z[c(1,3),2]
```

```
[1] "A" "A"
```

Calculate with matrices

```
> m <- rbind(c(1,2,3),  
             c(3,2,1),  
             c(5,6,4))  
  
> m  
[1,] 1 2 3  
[2,] 3 2 1  
[3,] 5 6 4
```

Calculate with matrices

```
> m <- rbind(c(1,2,3),  
             c(3,2,1),  
             c(5,6,4))
```

```
> m  
[1,] 1 2 3  
[2,] 3 2 1  
[3,] 5 6 4
```

```
> m + 1  
[1,] 2 3 4  
[2,] 4 3 2  
[3,] 6 7 5
```

Calculate with matrices

```
> m <- rbind(c(1,2,3),  
             c(3,2,1),  
             c(5,6,4))
```

```
> m  
[1,] 1 2 3  
[2,] 3 2 1  
[3,] 5 6 4
```

```
> m + 1  
[1,] 2 3 4  
[2,] 4 3 2  
[3,] 6 7 5
```

```
> m %*% solve(m) ## multiple m by its inverse
```

```
      [,1]      [,2]      [,3]  
[1,] 1 -2.220446e-16 2.220446e-16  
[2,] 0 1.000000e+00 2.220446e-16  
[3,] 0 0.000000e+00 1.000000e+00
```

Store and manipulate datasets: lists

A **list** is like a vector, but it can contain elements of different types.

```
> y <- diag(2)
> w <- list(x=1, y=y, z="abc")
```

Store and manipulate datasets: lists

A **list** is like a vector, but it can contain elements of different types.

```
> y <- diag(2)
> w <- list(x=1, y=y, z="abc")
```

```
w$y
```


Store and manipulate datasets: lists

A **list** is like a vector, but it can contain elements of different types.

```
> y <- diag(2)
> w <- list(x=1, y=y, z="abc")
```

```
w$y
```

```
      [,1] [,2]
[1,]    1    0
[2,]    0    1
```

```
names(r)
```

Store and manipulate datasets: lists

A **list** is like a vector, but it can contain elements of different types.

```
> y <- diag(2)
> w <- list(x=1, y=y, z="abc")
```

```
w$y
```

```
      [,1] [,2]
[1,]     1     0
[2,]     0     1
```

```
names(r)
```

```
[1] "x" "y" "z"
```

Store and manipulate datasets: data frames

A **data frame** is like a matrix but the columns can be different types. You could think of it as a list of vectors each of the same length.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> d <- data.frame(a=x, b=y)
```

Store and manipulate datasets: data frames

A **data frame** is like a matrix but the columns can be different types. You could think of it as a list of vectors each of the same length.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> d <- data.frame(a=x, b=y)
```

```
> d[2,]
  a b
2 1 B
```

Store and manipulate datasets: data frames

A **data frame** is like a matrix but the columns can be different types. You could think of it as a list of vectors each of the same length.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> d <- data.frame(a=x, b=y)
```

```
> d[2,]
  a b
2 1 B
```

```
> d$a
[1] 1 1 1 1 1
```

Store and manipulate datasets: data frames

A **data frame** is like a matrix but the columns can be different types. You could think of it as a list of vectors each of the same length.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> d <- data.frame(a=x, b=y)
```

```
> d[2,]
  a b
2 1 B
```

```
> d$a
[1] 1 1 1 1 1
```

```
> class(d$a)
[1] "numeric"
```

Store and manipulate datasets: data frames

A **data frame** is like a matrix but the columns can be different types. You could think of it as a list of vectors each of the same length.

```
> x <- rep(1, 5)
> y <- c("A", "B", "A", "B", "C")
> d <- data.frame(a=x, b=y)
```

```
> d[2,]
  a b
2 1 B
```

```
> d$a
[1] 1 1 1 1 1
```

```
> class(d$a)
[1] "numeric"
```

```
> class(d$b)
[1] "character"
```

`stop("Time for a break")`

`We will resume after Sys.sleep(10*60).`



Handle missing data

```
> y <- c("A", "B", "A", "B", "C")  
> y[2] <- 2  
> y[3] <- NA  
> y  
[1] "A" "2" NA "A" "B" "C"
```

Handle missing data

```
> y <- c("A", "B", "A", "B", "C")  
> y[2] <- 2  
> y[3] <- NA  
> y  
[1] "A" "2" NA "A" "B" "C"
```

```
> is.na(y)  
[1] FALSE FALSE TRUE FALSE FALSE FALSE
```

Handle missing data

```
> y <- c("A", "B", "A", "B", "C")  
> y[2] <- 2  
> y[3] <- NA  
> y  
[1] "A" "2" NA "A" "B" "C"
```

```
> is.na(y)  
[1] FALSE FALSE TRUE FALSE FALSE FALSE
```

```
> is.na(y[3])  
[1] TRUE
```

Handle missing data

```
> y <- c("A", "B", "A", "B", "C")  
> y[2] <- 2  
> y[3] <- NA  
> y  
[1] "A" "2" NA "A" "B" "C"
```

```
> is.na(y)  
[1] FALSE FALSE TRUE FALSE FALSE FALSE
```

```
> is.na(y[3])  
[1] TRUE
```

```
> na.omit(y)  
[1] "A" "2" "A" "B" "C"
```

Save and load data: working directory

To open files, you will need to tell R where the files are. To do this, it helps to know where R will start looking. This is called the 'working directory'.

```
> getwd()  
[1] "C:/"
```

Save and load data: working directory

To open files, you will need to tell R where the files are. To do this, it helps to know where R will start looking. This is called the 'working directory'.

```
> getwd()  
[1] "C:/"
```

You can change this using `setwd`.

```
> setwd("0:/Documents")
```

The working directory can also be chosen using the Session menu in RStudio.

Save and load data: reading csv files

Suppose I have a spreadsheet stored in CSV format.

```
"id","age","sex","diet","bmi"  
1,32,"M",0,25.7957231474661  
2,35,"M",0,28.8952139451377  
3,41,"M",0,29.9258448186199  
4,29,"M",0,27.3383500990741  
5,33.5,"M",1,28.2469210985821  
6,33.2,"M",1,27.0473176810354  
7,32.9,"M",1,30.3031786852156  
8,32.6,"F",1,28.5621419729205  
9,32.3,"F",1,28.05344654591  
10,32,"F",1,28.8864676350323  
...
```

Save and load data: reading csv files

Suppose I have a spreadsheet stored in CSV format.

`read.csv` will read csv files and save them as data frames.

```
"id","age","sex","diet","bmi"  
1,32,"M",0,25.7957231474661  
2,35,"M",0,28.8952139451377  
3,41,"M",0,29.9258448186199  
4,29,"M",0,27.3383500990741  
5,33.5,"M",1,28.2469210985821  
6,33.2,"M",1,27.0473176810354  
7,32.9,"M",1,30.3031786852156  
8,32.6,"F",1,28.5621419729205  
9,32.3,"F",1,28.05344654591  
10,32,"F",1,28.8864676350323  
...
```

```
> dat <- read.csv("bmi.csv")
```

```
> dat[1:5,]  
  id age sex diet    bmi  
1  1 32.0  M    0 25.79572  
2  2 35.0  M    0 28.89521  
3  3 41.0  M    0 29.92584  
4  4 29.0  M    0 27.33835  
5  5 33.5  M    1 28.24692
```

```
> dat$bmi[1:5]  
[1] 25.79572 28.89521 29.92584 27.33835 28.24692  
> dat[5,]  
  id age sex diet    bmi  
5  5 33.5  M    1 28.24692
```


Save and load data: writing csv files

`write.csv` saves a data frame as a csv file.

```
> dat.corrected <- dat  
> dat.corrected$sex[5] <- "F" ## make correction  
> write.csv(dat.corrected, "bmi-corrected.csv", row.names=F, quote=F)
```

`write.table` is similar but allows the user to change the column separator character. Here we separate columns by a semicolon rather than a comma.

```
> write.table(dat, "bmi-corrected.csv", row.names=F, quote=F, sep=";")
```

Save and load data: writing csv files

`write.csv` saves a data frame as a csv file.

```
> dat.corrected <- dat
> dat.corrected$sex[5] <- "F" ## make correction
> write.csv(dat.corrected, "bmi-corrected.csv", row.names=F, quote=F)
```

`write.table` is similar but allows the user to change the column separator character. Here we separate columns by a semicolon rather than a comma.

```
> write.table(dat, "bmi-corrected.csv", row.names=F, quote=F, sep=";")
```

Note: There is similarly a function `read.table()` just like `read.csv()` but more flexible.

Statistical analyses: basic summaries

```
> mean(dat$bmi)      ## mean  
[1] 27.60761
```

Statistical analyses: basic summaries

```
> mean(dat$bmi)          ## mean  
[1] 27.60761
```

```
> median(dat$bmi)        ## median  
[1] 27.58069
```

Statistical analyses: basic summaries

```
> mean(dat$bmi)          ## mean  
[1] 27.60761
```

```
> median(dat$bmi)        ## median  
[1] 27.58069
```

```
> sd(dat$bmi)             ## standard deviation  
[1] 1.47353
```

Statistical analyses: basic summaries

```
> mean(dat$bmi)          ## mean  
[1] 27.60761
```

```
> median(dat$bmi)        ## median  
[1] 27.58069
```

```
> sd(dat$bmi)             ## standard deviation  
[1] 1.47353
```

```
> min(dat$bmi)            ## min  
[1] 25.34356
```

Statistical analyses: basic summaries

```
> mean(dat$bmi)          ## mean  
[1] 27.60761
```

```
> median(dat$bmi)        ## median  
[1] 27.58069
```

```
> sd(dat$bmi)             ## standard deviation  
[1] 1.47353
```

```
> min(dat$bmi)            ## min  
[1] 25.34356
```

```
> quantile(dat$bmi, probs=0.25) ## first quartile  
25%  
26.60918
```

Statistical analyses: basic summaries

```
> mean(dat$bmi)          ## mean  
[1] 27.60761
```

```
> median(dat$bmi)        ## median  
[1] 27.58069
```

```
> sd(dat$bmi)            ## standard deviation  
[1] 1.47353
```

```
> min(dat$bmi)           ## min  
[1] 25.34356
```

```
> quantile(dat$bmi, probs=0.25) ## first quartile  
      25%  
26.60918
```

```
> table(dat$sex)          ## frequencies  
F  M  
13 7
```


Statistical analyses: dataset summaries

The `summary` function can be used to summarize single variables

```
> summary(dat$bmi)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 25.34  26.61   27.58   27.61  28.64   30.30
```

Statistical analyses: dataset summaries

The summary function can be used to summarize single variables

```
> summary(dat$bmi)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 25.34  26.61   27.58   27.61  28.64   30.30
```

or entire datasets.

```
> summary(dat)
      id          age          sex          diet          bmi
Min.   : 1.00   Min.   :29.00   Length:20   Min.     :0.00   Min.     :25.34
1st Qu.: 5.75   1st Qu.:30.12   Class :character 1st Qu.   :0.00   1st Qu.  :26.61
Median :10.50   Median :31.55   Mode  :character Median    :1.00   Median   :27.58
Mean   :10.50   Mean   :31.85                Mean     :0.55   Mean     :27.61
3rd Qu.:15.25   3rd Qu.:32.67                3rd Qu.  :1.00   3rd Qu.  :28.64
Max.   :20.00   Max.   :41.00                Max.     :1.00   Max.     :30.30
```

Statistical analyses: evaluating associations

There is a fairly strong association between BMI and age.

```
> cor(dat$bmi, dat$age)
[1] 0.5705649
```

Here is the association.

```
> fit <- lm(bmi~age, data=dat)
> fit
Call:
lm(formula = bmi ~ age, data = dat)

Coefficients:
(Intercept)          age
    17.6770         0.3118
```

Statistical analyses: evaluating associations

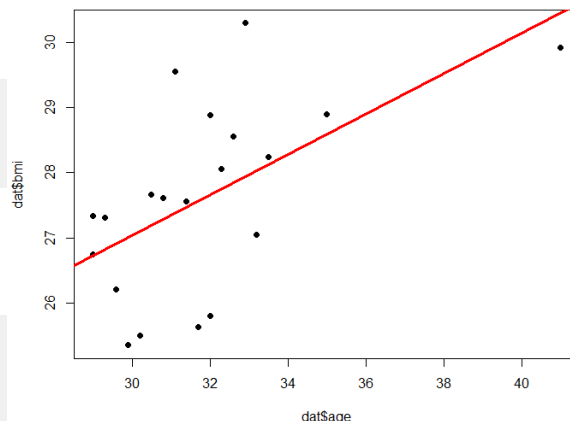
There is a fairly strong association between BMI and age.

```
> cor(dat$bmi, dat$age)
[1] 0.5705649
```

Here is the association.

```
> fit <- lm(bmi~age, data=dat)
> fit
Call:
lm(formula = bmi ~ age, data = dat)
```

```
Coefficients:
(Intercept)      age
  17.6770      0.3118
```



The **regression line** in the plot runs from $(29, 17.6770 + 0.3118 \times 29) = (29, 26.7192)$ to $(41, 17.6770 + 0.3118 \times 41) = (41, 30.4608)$.

Statistical analyses: summarizing regression model fits

```
> summary(fit)
```

Call:

```
lm(formula = bmi ~ age, data = dat)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|--------|--------|--------|
| -1.9334 | -0.7760 | 0.2152 | 0.5293 | 2.3682 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) | |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 17.6770 | 3.3805 | 5.229 | 5.67e-05 | *** |
| age | 0.3118 | 0.1058 | 2.948 | 0.00861 | ** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.243 on 18 degrees of freedom

Multiple R-squared: 0.3255, Adjusted R-squared: 0.2881

F-statistic: 8.688 on 1 and 18 DF, p-value: 0.008612

Statistical analyses: fitting multiple variable models

```
> fit <- lm(bmi ~ age + sex + diet, data=dat)
```

Statistical analyses: fitting multiple variable models

```
> fit <- lm(bmi ~ age + sex + diet, data=dat)
```

```
> summary(fit)
```

Call:

```
lm(formula = bmi ~ age + sex + diet, data = dat)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|--------|--------|--------|
| -2.2978 | -0.6210 | 0.0106 | 0.8208 | 1.7891 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) | |
|-------------|----------|------------|---------|----------|-----|
| (Intercept) | 18.4784 | 3.7464 | 4.932 | 0.00015 | *** |
| age | 0.2643 | 0.1228 | 2.152 | 0.04700 | * |
| sexM | 0.3512 | 0.6866 | 0.512 | 0.61596 | |
| diet | 1.0671 | 0.5532 | 1.929 | 0.07168 | . |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.188 on 16 degrees of freedom

Multiple R-squared: 0.4528, Adjusted R-squared: 0.3502

F-statistic: 4.413 on 3 and 16 DF, p-value: 0.01921

Statistical analyses: retrieving model fit elements

```
> coef(fit)
(Intercept)      age      sexM      diet
 18.4783563    0.2643460    0.3512052    1.0671192
```


Statistical analyses: retrieving model fit elements

```
> coef(fit)
(Intercept)      age      sexM      diet
 18.4783563    0.2643460    0.3512052    1.0671192
```

```
> coef(summary(fit))
              Estimate Std. Error  t value    Pr(>|t|)
(Intercept) 18.4783563   3.7463529  4.9323587 0.0001500352
age          0.2643460   0.1228288  2.1521491 0.0469973143
sexM         0.3512052   0.6865686  0.5115369 0.6159621270
diet         1.0671192   0.5532367  1.9288657 0.0716810178
```

Statistical analyses: retrieving model fit elements

```
> coef(fit)
(Intercept)      age      sexM      diet
 18.4783563    0.2643460    0.3512052    1.0671192
```

```
> coef(summary(fit))
              Estimate Std. Error  t value    Pr(>|t|)
(Intercept) 18.4783563   3.7463529  4.9323587 0.0001500352
age          0.2643460   0.1228288  2.1521491 0.0469973143
sexM         0.3512052   0.6865686  0.5115369 0.6159621270
diet         1.0671192   0.5532367  1.9288657 0.0716810178
```

```
> coef(summary(fit))["age", "Estimate"]
[1] 0.264346
```

Statistical analyses: retrieving model fit elements

```
> coef(fit)
(Intercept)      age      sexM      diet
 18.4783563    0.2643460    0.3512052    1.0671192
```

```
> coef(summary(fit))
              Estimate Std. Error  t value    Pr(>|t|)
(Intercept) 18.4783563   3.7463529  4.9323587 0.0001500352
age          0.2643460   0.1228288  2.1521491 0.0469973143
sexM         0.3512052   0.6865686  0.5115369 0.6159621270
diet         1.0671192   0.5532367  1.9288657 0.0716810178
```

```
> coef(summary(fit))["age", "Estimate"]
[1] 0.264346
```

```
> names(summary(fit))
[1] "call"      "terms"      "residuals"
[4] "coefficients" "aliased"    "sigma"
[7] "df"        "r.squared"  "adj.r.squared"
[10] "fstatistic" "cov.unscaled"
```

Statistical analyses: retrieving model fit elements

```
> coef(fit)
(Intercept)      age      sexM      diet
 18.4783563    0.2643460    0.3512052    1.0671192
```

```
> coef(summary(fit))
              Estimate Std. Error  t value    Pr(>|t|)
(Intercept) 18.4783563   3.7463529  4.9323587 0.0001500352
age          0.2643460   0.1228288  2.1521491 0.0469973143
sexM         0.3512052   0.6865686  0.5115369 0.6159621270
diet         1.0671192   0.5532367  1.9288657 0.0716810178
```

```
> coef(summary(fit))["age", "Estimate"]
[1] 0.264346
```

```
> names(summary(fit))
[1] "call"      "terms"      "residuals"
[4] "coefficients" "aliases"    "sigma"
[7] "df"        "r.squared"  "adj.r.squared"
[10] "fstatistic" "cov.unscaled"
```

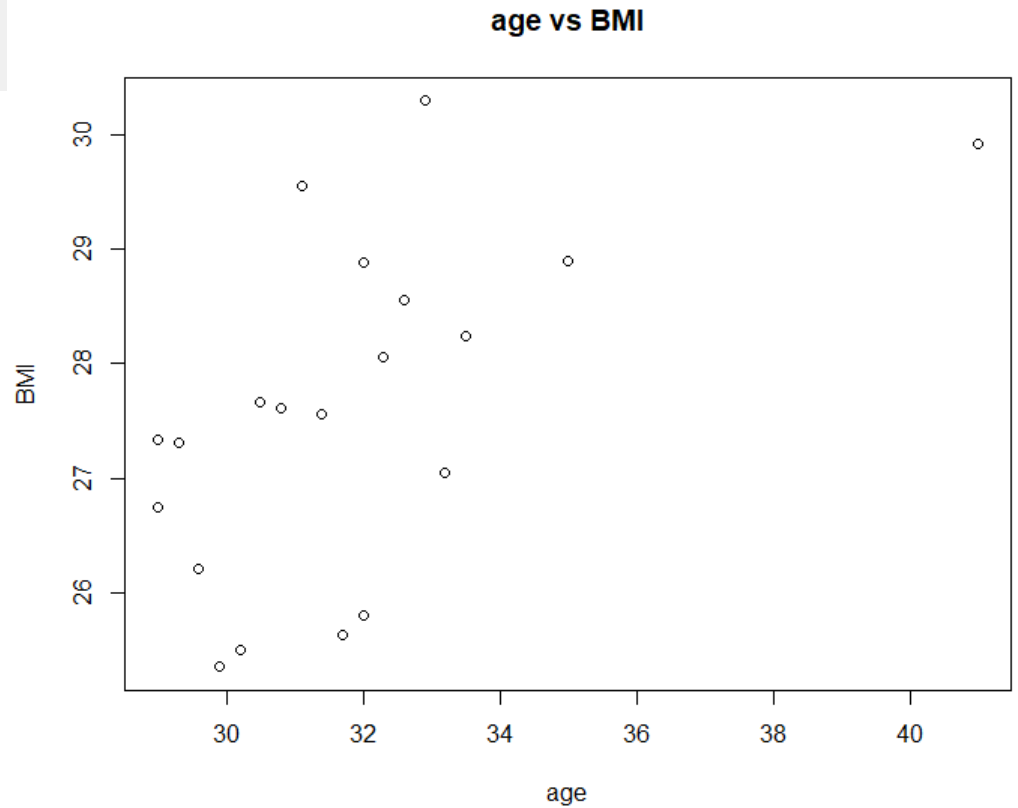
```
> summary(fit)$adj.r.squared
[1] 0.3502004
```

Visualise: scatterplot

```
plot(dat$age, dat$bmi,  
     main="age vs BMI",  
     xlab="age",  
     ylab="BMI")
```

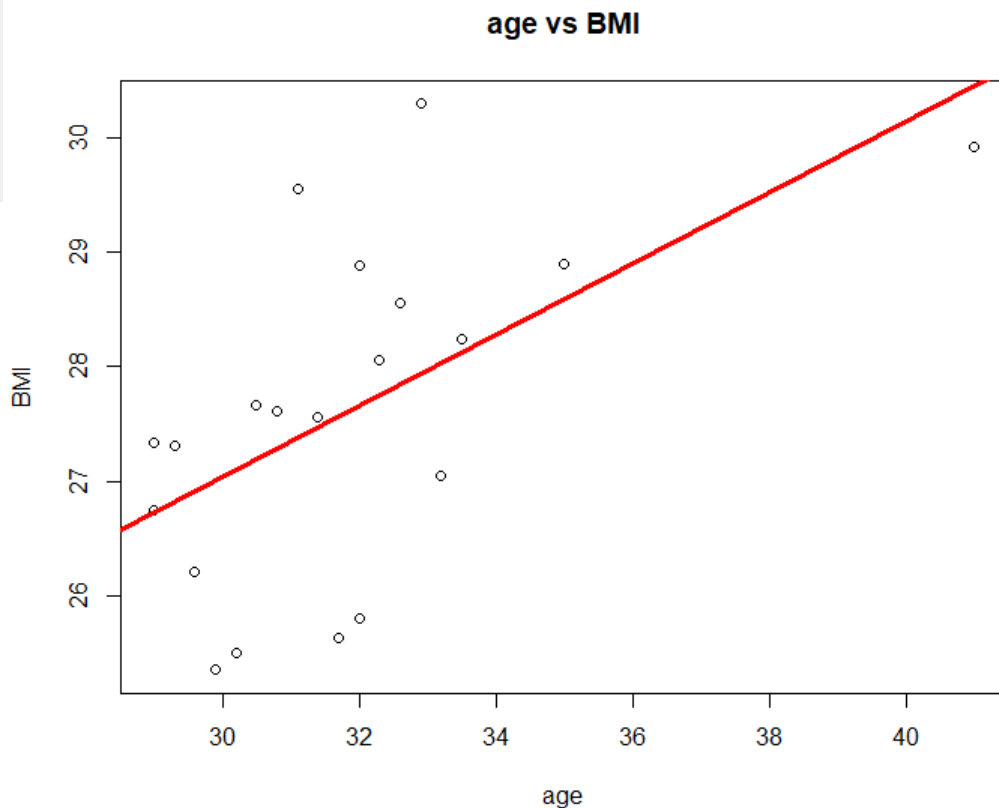
Visualise: scatterplot

```
plot(dat$age, dat$bmi,  
     main="age vs BMI",  
     xlab="age",  
     ylab="BMI")
```



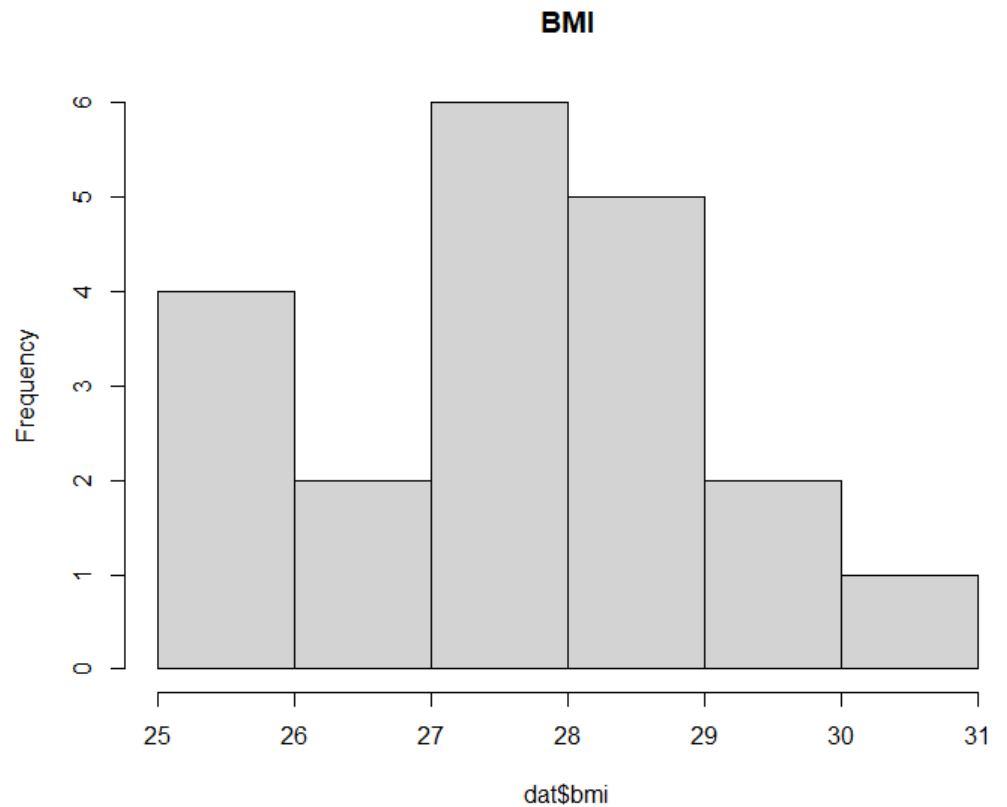
Visualise: scatterplot with regression line

```
plot(dat$age, dat$bmi,  
     main="age vs BMI",  
     xlab="age",  
     ylab="BMI")  
fit <- lm(bmi ~ age,  
          data=dat)  
abline(fit, col="red", lwd=3)
```



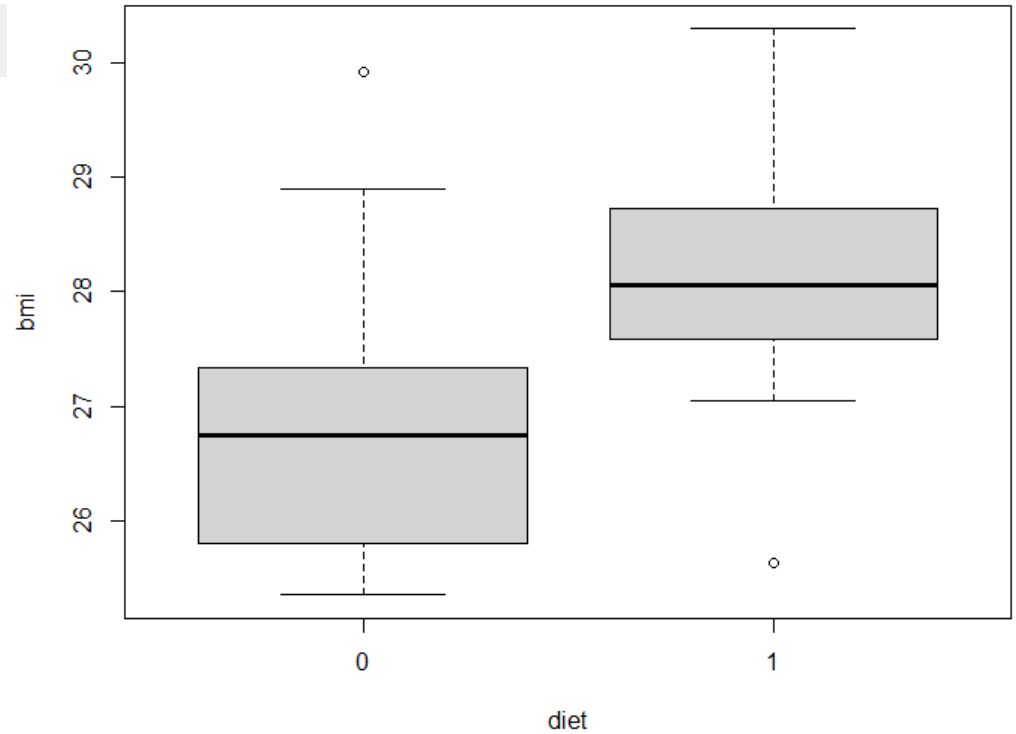
Visualise: histogram

```
hist(dat$bmi, main="BMI")
```

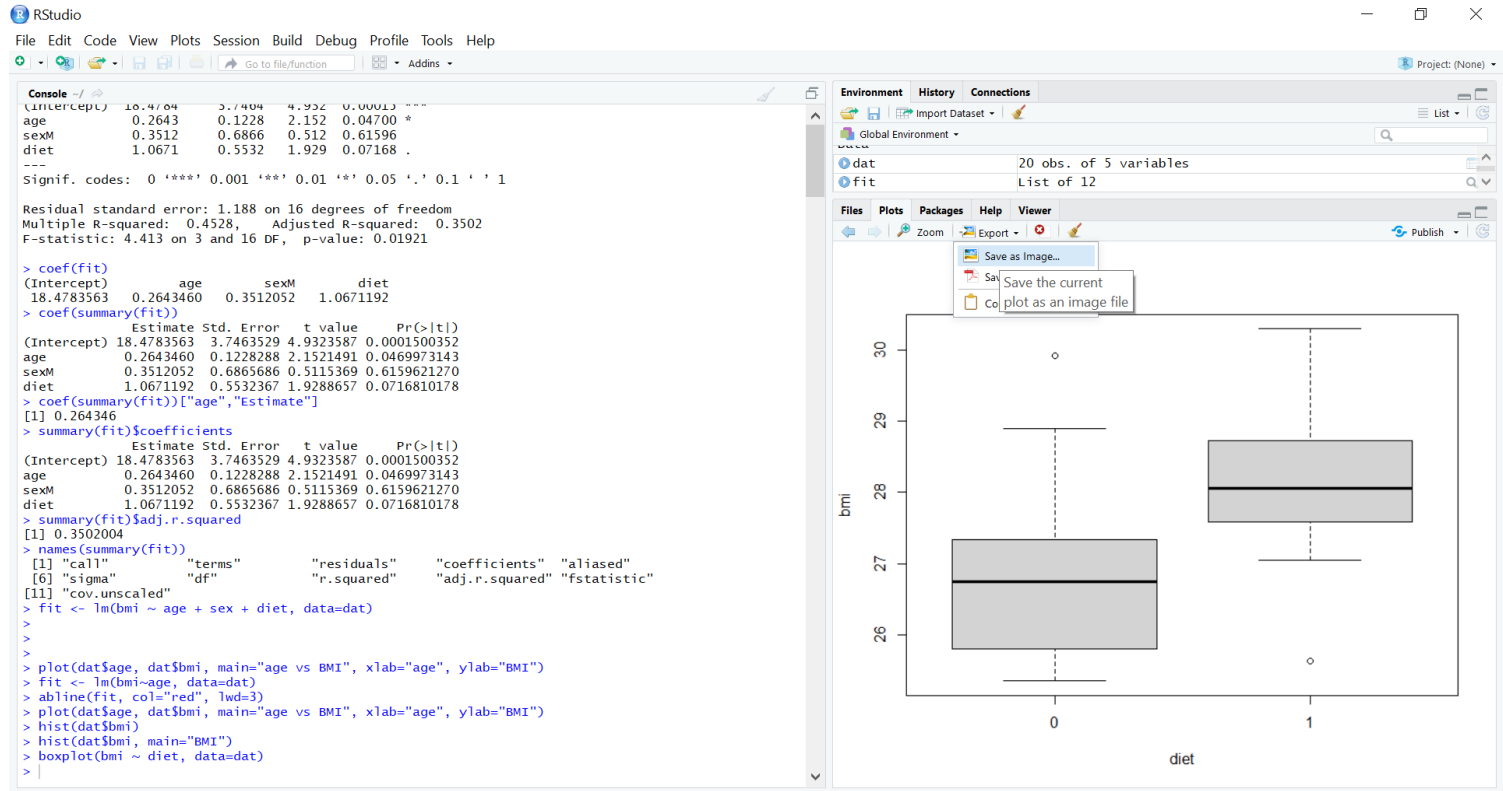


Visualise: boxplot

```
boxplot(bmi ~ diet, data=dat)
```



Visualise: saving



Visualise: saving *like a total maniac*

```
> png(filename="boxplot.png", width=10, height=10, units="cm", res=500)
> boxplot(bmi~diet, data=dat)
> dev.off()
```

Recipes: create and apply

Functions consist of a sequence of commands applied to a set of variables that return some output.

R provides many functions such as `help`, `cor`, `median`.

Recipes: create and apply

Functions consist of a sequence of commands applied to a set of variables that return some output.

R provides many functions such as `help`, `cor`, `median`.

Users can create their own functions for repetitive tasks, e.g. a function for the distance of a point from 0.

```
euclidean.norm <- function(x) {  
  res <- sqrt(sum(x^2))  
  return(res)  
}
```

Recipes: create and apply

Functions consist of a sequence of commands applied to a set of variables that return some output.

R provides many functions such as `help`, `cor`, `median`.

Users can create their own functions for repetitive tasks, e.g. a function for the distance of a point from 0.

```
euclidean.norm <- function(x) {  
  res <- sqrt(sum(x^2))  
  return(res)  
}
```

```
> euclidean.norm(c(3,4))  
[1] 5  
  
> euclidean.norm(c(2,3,6))  
[1] 7  
  
> euclidean.norm(c(1,4,8))  
[1] 9
```

Recipes: repeat

Given a vector people:

```
> people <- c("bill",  
              "jane",  
              "bob",  
              "felicia",  
              "carl",  
              "apple")  
> scores <- c(90, 80,  
              60, 95,  
              75, 99)
```

Suppose we want to
make a list of people
who obtained high
scores.

Recipes: repeat

Given a vector people: **Step 1.** Create a function for determining a high score.

```
> people <- c("bill",  
              "jane",  
              "bob",  
              "felicia",  
              "carl",  
              "apple")  
> scores <- c(90, 80,  
              60, 95,  
              75, 99)
```

```
is.high <- function(s) {  
  return(s > 90)  
  ## or something  
  ## really complicated  
  ## takes into account  
  ## grade barriers  
  ## and homework  
  ## completed and  
  ## the name of  
  ## first pet.  
}
```

Suppose we want to make a list of people who obtained high scores.

Recipes: repeat

Given a vector people: **Step 1.** Create a function for determining a high score.

```
> people <- c("bill",  
             "jane",  
             "bob",  
             "felicia",  
             "carl",  
             "apple")  
> scores <- c(90, 80,  
             60, 95,  
             75, 99)
```

Suppose we want to make a list of people who obtained high scores.

```
is.high <- function(s) {  
  return(s > 90)  
  ## or something  
  ## really complicated  
  ## takes into account  
  ## grade barriers  
  ## and homework  
  ## completed and  
  ## the name of  
  ## first pet.  
}
```

Step 2. Construct the list of people by applying is.high().

Two equivalent ways:

Option 1. For-loop

```
n <- length(people)  
high <- character()  
for (i in 1:n) {  
  if (is.high(scores[i]))  
    high <- c(high, people[i])  
}
```

Option 2. sapply()

```
yes <- sapply(scores, is.high)  
high <- people[yes]
```

Decisions: ask questions

```
> x <- 11:15
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13          ## 13 equals 13  
[1] TRUE
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13          ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12          ## 13 not equal to 12
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13          ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12          ## 13 not equal to 12
```

```
> x[3] < x[4]          ## 13 less than 14
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13      ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12      ## 13 not equal to 12
```

```
> x[3] < x[4]     ## 13 less than 14
```

```
> x[1] <= x[2] & !x[2] > x[3]  ## 11 <= 12 and 12 not greater than 13
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13      ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12      ## 13 not equal to 12
```

```
> x[3] < x[4]     ## 13 less than 14
```

```
> x[1] <= x[2] & !x[2] > x[3]  ## 11 <= 12 and 12 not greater than 13
```

```
> x[1] < x[2] | x[2] > x[3]    ## 11 < 12 or 12 > 13
```

Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13          ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12          ## 13 not equal to 12
```

```
> x[3] < x[4]          ## 13 less than 14
```

```
> x[1] <= x[2] & !x[2] > x[3]  ## 11 <= 12 and 12 not greater than 13
```

```
> x[1] < x[2] | x[2] > x[3]    ## 11 < 12 or 12 > 13
```

```
> x > 12  
[1] TRUE FALSE FALSE FALSE FALSE
```


Decisions: ask questions

```
> x <- 11:15
```

```
> x[3] == 13          ## 13 equals 13  
[1] TRUE
```

```
> x[3] != 12          ## 13 not equal to 12
```

```
> x[3] < x[4]          ## 13 less than 14
```

```
> x[1] <= x[2] & !x[2] > x[3]  ## 11 <= 12 and 12 not greater than 13
```

```
> x[1] < x[2] | x[2] > x[3]    ## 11 < 12 or 12 > 13
```

```
> x > 12  
[1] TRUE FALSE FALSE FALSE FALSE
```

```
> all(x < 12)          ## each value in x < 12  
[1] TRUE
```

Decisions: alternatives

```
if (score > 85) {  
  grade <- "A"  
} else {  
  grade <- "F"  
}
```

Decisions: alternatives

```
if (score > 85) {  
  grade <- "A"  
} else {  
  grade <- "F"  
}
```

```
grade <- ifelse(score > 85, "A", "F")
```

Decisions: alternatives

```
if (score > 85) {  
  grade <- "A"  
} else {  
  grade <- "F"  
}
```

```
grade <- ifelse(score > 85, "A", "F")
```

```
if (score > 85) {  
  grade <- "A"  
} else if (score > 75) {  
  grade <- "B"  
} else {  
  grade <- "F"  
}
```

Decisions: alternatives

```
if (score > 85) {  
  grade <- "A"  
} else {  
  grade <- "F"  
}
```

```
grade <- ifelse(score > 85, "A", "F")
```

```
if (score > 85) {  
  grade <- "A"  
} else if (score > 75) {  
  grade <- "B"  
} else {  
  grade <- "F"  
}
```

```
grade <- ifelse(score > 85, "A", ifelse(score > 75, "B", "F"))
```

Share: loading packages

An **R package** is a collection of functions and/or data sets created for use by other R users. A package can be loaded using `library`.

```
> library(Hmisc)  
Attaching package: 'Hmisc'
```

Share: loading packages

An **R package** is a collection of functions and/or data sets created for use by other R users. A package can be loaded using `library`.

```
> library(Hmisc)
Attaching package: 'Hmisc'
```

After loading, functions and data provided by the package can be used.

```
> describe(c(1,1,2,1,3,5,1,3,5,2,3,4))
      n missing distinct      Info
  12      0         5      0.944
  Mean      Gmd
 2.583     1.742

lowest : 1 2 3 4 5, highest: 1 2 3 4 5

Value      1      2      3      4
Frequency   4      2      3      1
Proportion 0.333 0.167 0.250 0.083

Value      5
Frequency   2
Proportion 0.167
```

Share: installing packages

If the package has not been installed, you'll see an error message like this.

```
> library(Hmisc)  
Error in library(Hmisc) : there is no package called 'Hmisc'
```


Share: installing packages

If the package has not been installed, you'll see an error message like this.

```
> library(Hmisc)
Error in library(Hmisc) : there is no package called 'Hmisc'
```

The package can be installed using `install.packages`.

```
> install.packages("Hmisc")

Installing package into 'C:/Users/dude/Documents/R/win-library/4.0'
(as 'lib' is unspecified)
also installing the dependencies 'png', 'jpeg', 'Formula',
'latticeExtra', 'gridExtra', 'htmlTable', 'viridis'

trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.0/png_0.1-7.zip'
Content type 'application/zip' length 336667 bytes (328 KB)
downloaded 328 KB

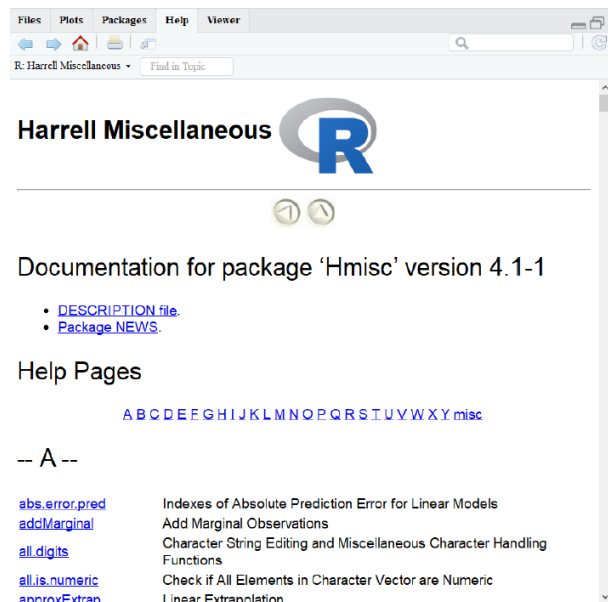
...

package 'Hmisc' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
C:\Users\dude\AppData\Local\Temp\RtmpS221DR\downloaded_packages
```

Share: package help

The *help* function can be used to open the documentation for an R package, e.g. `help(package="Hmisc")`



Search the package list on CRAN
<https://cran.r-project.org/web/packages/Hmisc/>

Hmisc: Harrell Miscellaneous

Contains many functions useful for data analysis, high-level graphics, utility operations, advanced table making, variable clustering, character string manipulation, etc.

Version: 4.4-1
Depends: [lattice](#), [survival](#) (≥ 3.1-6), [Formula](#), [ggplot2](#) (≥ 2.2)
Imports: methods, [latticeExtra](#), [cluster](#), [rpart](#), [nnet](#), [foreign](#), [gtable](#), grid, [g](#)
Suggests: [acepack](#), [chron](#), [rms](#), [mice](#), [tables](#), [knitr](#), [plotly](#) (≥ 4.5.6), [rlang](#), [r](#)
Published: 2020-08-10
Author: Frank E Harrell Jr, with contributions from Charles Dupont and
Maintainer: Frank E Harrell Jr <fh at fharrell.com>
License: [GPL-2](#) | [GPL-3](#) [expanded from: GPL (≥ 2)]
URL: <https://hbiostat.org/R/Hmisc/>, <https://github.com/harrelfe/Hmisc>
NeedsCompilation: yes
Materials: [README](#) [NEWS](#) [ChangeLog](#)
In views: [Bayesian](#), [ClinicalTrials](#), [Econometrics](#), [MissingData](#), [Multivari](#)
CRAN checks: [Hmisc results](#)

Downloads:

Reference manual: [Hmisc.pdf](#)
Package source: [Hmisc_4.4-1.tar.gz](#)
Windows binaries: r-devel: [Hmisc_4.4-1.zip](#), r-release: [Hmisc_4.4-1.zip](#), r-oldrel: [H](#)

Share: learn R with "swirl"

swirl (<http://swirlstats.com/students.html>) allows you to learn R within R itself.



```
> install.packages("swirl")  
> library(swirl)  
| Hi! Type swirl() when you are ready to begin.
```

Share: creating packages for others

If you have a set of R functions and/or a dataset that you think others might like to use, create a package and put them in it!

Here is a good place to get started:

<https://support.rstudio.com/hc/en-us/articles/200486488-Developing-Packages-with-RStudio>

Acknowledgements

Slides were based on work by these beautiful minds.



Harriet Mills

University of Bristol



Andrew Simpkin

National University of
Ireland



James Staley

UCB